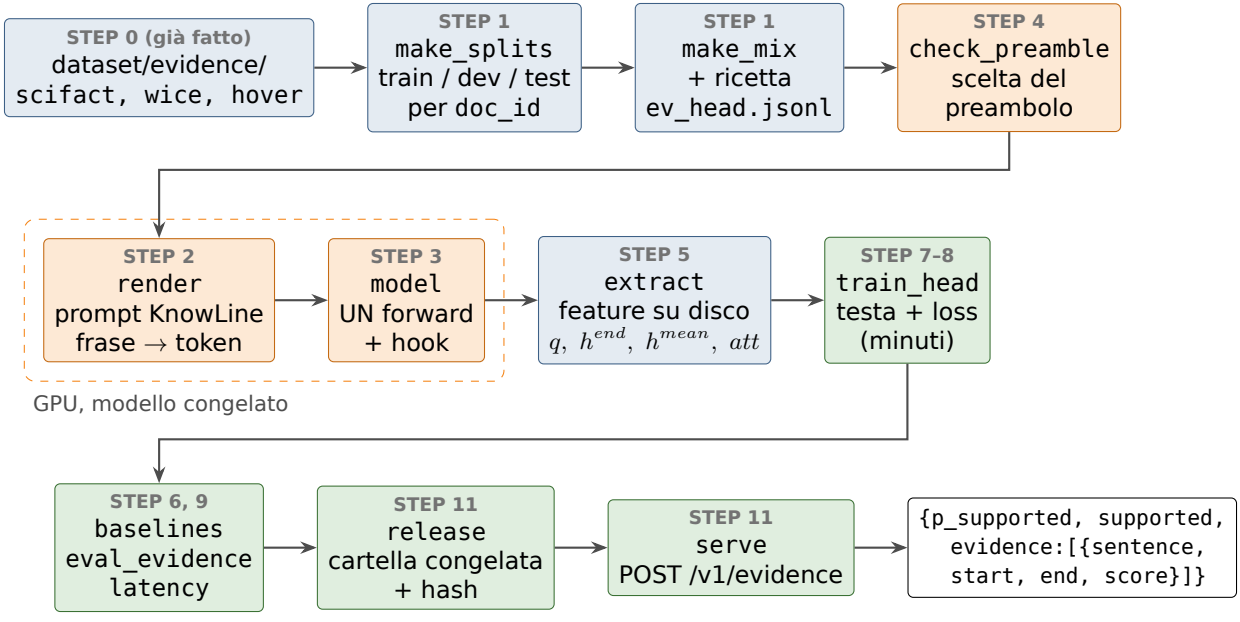
STEP 0 (già fatto)
dataset/evidence/
scifact, wice, hover
STEP 1
make_splits
train / dev / test
per doc_id
STEP 1
make_mix
+ ricetta
ev_head.jsonl
STEP 4
check_preamble
scelta del
preambolo
STEP 2
render
prompt KnowLine
frase → token
STEP 3
model
UN forward
+ hook
STEP 5
extract
feature su disco
$q,\ h^{end},\ h^{mean},\ att$
STEP 7-8
train_head
testa + loss
(minuti)
GPU, modello congelato
STEP 6, 9
baselines
eval_evidence
latency
STEP 11
release
cartella congelata
+ hash
STEP 11
serve
POST /v1/evidence
{p_supported, supported,
evidence:[{sentence,
start, end, score}]}